

Lexing

Goal

- Input: source code as a character sequence

`a = 3 + sin(mx)`

- Goal: split source code into lexical structures (*lexemes*)

`a , = , 3 , + , sin , (, mx ,)`

- Output: a sequence of syntactic atoms (*symbols*) represented by those lexemes

- Often implemented as: Symbol = <Token, Attribute>

- For our example, we get this sequence of symbols:

<Identifier, "a">

<AssignmentOperator,>

<Number, 3>

<ArithmeticOperator, "+"> (or just <AddOperator,>)

<Identifier, "sin">

<SpecialCharacter, "("> (or just <OpenParenthesis,>)

<Identifier, "mx">

<SpecialCharacter, ")"> (or just <CloseParenthesis,>)

Lexical Analysis

- There are compilers without a separate parser and lexer
 - The parser could directly read and process the source code
- In practice, it can be useful to have a lexer
 - Simplifies the job of the parser. The lexer can do a first preprocessing of the source code, for example filter out whitespace characters, newlines, comments,...
 - The lexer is less complex than the parser and can be implemented very efficiently
- Note. You don't lex the entire source code before starting the parser.
- Typical implementation:

```
class Lexer {  
    Lexer(Reader sourcefile) { ... }  
    Symbol getNextSymbol() { ... }  
}
```

Defining lexemes

- We can use a *regular expression* to describe the allowed lexemes in the source code

`"=" | "+" | "-" | "++" | "(" | ")" | [a-zA-Z] [a-zA-Z]* | ...`

Choice

Short for `"+" "+"`

Short for: `"a" | "b" | ...`

Zero or more repetitions
(Kleene star)

- For a lexer, it is usually more convenient to write one RE per symbol class:

AssignmentOperator: `"="`

IncrementOperator: `"++"`

ArithmeticOperator: `"+" | "-"`

SpecialCharacter: `"(" | ")"`

Identifier: `[a-zA-Z] [a-zA-Z]*`

Number: `[1-9] [0-9]*`

Whitespace: `(" " | "\t" | "\n")*`

Formal Definition of Regular Expressions (RE)

- Given: an alphabet Ω of allowed characters $a, b, \dots \in \Omega$
- The set of regular expressions RE over Ω is defined as:
 - $\varepsilon \in RE$ (empty expression)
 - $\Omega \subseteq RE$
 - If $\alpha, \beta \in RE$ then also
 - $\alpha \beta \in RE$
 - $\alpha \mid \beta \in RE$
 - $\alpha^* \in RE$

Regular languages

- A regular expression specifies a *regular language*, i.e., the possible sequences of characters described by that regular expression.
- The mapping $\llbracket \cdot \rrbracket: RE \rightarrow 2^{\Omega^*}$ gives the *semantics* of a regular expression
 - $\llbracket \varepsilon \rrbracket = \{\}$ (empty language)
 - $\llbracket a \rrbracket = \{a\}$
 - $\llbracket \alpha \beta \rrbracket = \llbracket \alpha \rrbracket \llbracket \beta \rrbracket$
 - $\llbracket \alpha | \beta \rrbracket = \llbracket \alpha \rrbracket \cup \llbracket \beta \rrbracket$
 - $\llbracket \alpha^* \rrbracket = \llbracket \alpha \rrbracket^*$
- Example: The RE
 $[1-9] [0-9]^*$
specifies the regular language
 $\llbracket [1-9] [0-9]^* \rrbracket = \{ 1, 2, 3, \dots, 9, 10, 11, 12, \dots, 1340242, 4595983, \dots \}$


Simple Matching Problem

- Before we show how we can implement a lexer using REs, let's look at a simpler problem:
 - *I give you a sequence of characters and a RE and you have to decide whether the sequence is in the language of the RE*
 - Often used to validate user input, for example a phone number
- Example: Input: 1233072, RE: [1-9] [0-9]*
- Implementation:

```
boolean match(Reader r) {  
    char c = r.read(); // read returns -1 if end of input reached  
    if(c<'1' || c>'9') return false;  
    while(true) {  
        c = r.read();  
        if(c==-1) break; // end of input reached  
        if(c<'0' || c>'9') return false;  
    }  
    return true; // all good  
}
```

Nondeterministic Finite Automaton (NFA)


- Can any RE be translated to simple code? A more theoretical approach is needed for this
- Here is an example Nondeterministic Finite Automaton (NFA)



- An NFA consists of
 - Finite set S of states
 - Initial state $s_0 \in S$
 - Set $F \subseteq S$ of final states
 - Input alphabet Ω
 - Set of transitions $T: S \times \Omega \cup \{\varepsilon\} \times S$
- If the NFA is in state $s \in S$ and consumes a character c it moves to state $t \in S$ provided there is a transition $(s, c, t) \in T$

We write $s \xrightarrow{c} t$


Simple Matching Problem with NFA




- For the input $a_1 a_2 a_3 \dots a_n$ with $a_i \in \text{alphabet } \Omega$, we say that the NFA accepts $a_1 a_2 a_3 \dots a_n$ if there is a sequence of transitions from the initial state to a final state that consume $a_1 a_2 a_3 \dots a_n$
- Example: Input abbbc
 - Sequence of transitions: $0 \xrightarrow{a} 1 \xrightarrow{b} 1 \xrightarrow{b} 1 \xrightarrow{b} 1 \xrightarrow{c} 2$ Accepted!
- Example: Input abbba
 - Sequence of transitions: $0 \xrightarrow{a} 1 \xrightarrow{b} 1 \xrightarrow{b} 1 \xrightarrow{b} 1 \xrightarrow{a} \text{X}$ Not accepted!
- Congratulations! We have built an NFA that accepts the language of the RE $a b^* c$
- *Kleene's theorem*: the set of language recognized by NFAs is identical to the set of languages of REs, i.e., we can always construct an NFA for an RE

Choice

- A RE of the form $a \mid b$ can be represented by an NFA with multiple transitions leaving a state and ϵ -transitions
- ϵ -transitions are “empty” transition. They don’t consume characters.
- Example: NFA with five states for the RE $(a \mid b) c$




- Note that ϵ -transitions can be always removed:



Implementing the Simple Matching Problem

- Thanks to the representation of the RE by a Finite Automaton, we can easily write down an implementation (or even write a tool that automatically translates a RE to code):



```
boolean simpleMatch(Reader r) {  
    int state = 0; // initial state  
    while(true) {  
        char c = r.read();  
        if(c==-1) break; // end of input reached  
        if(state==0) {  
            if(c=='a') state = 1;  
            else return false; // input not accepted  
        }  
        else if(state==1) {  
            if(c=='b') state = 1;  
            else if(c=='c') state = 2;  
            else return false; // input not accepted  
        }  
        else return false;  
    }  
    return state==2; // in final state?  
}
```