

Languages & Translators

LINGI2132

JVM Bytecode Instructions

Nicolas LAURENT

Université catholique de Louvain

Java to Bytecode

```
2 public class Main {  
3     public int foo() {  
4         System.out.println("Hello, world!");  
5         return 42;  
6     }  
7 }
```

> javac Main.java

> javap -v -p Main.class

or use <https://javap.yawk.at>

▼ Main.class

```
1 Classfile /tmp/6369527420976196219/classes/Main.class  
2   Last modified Apr 7, 2021; size 470 bytes  
3   SHA-256 checksum cc6d07f8dabc50d0867300b378ac6b3e638938c7e9db576885c9dd7625931624  
4   Compiled from "Main.java"  
5 public class Main  
6   minor version: 0  
7   major version: 60  
8   flags: (0x0021) ACC_PUBLIC, ACC_SUPER  
9   this_class: #21 // Main  
10  super_class: #2 // java/lang/Object  
11  interfaces: 0, fields: 0, methods: 2, attributes: 1
```

Constant Pool

```
12 Constant pool:
13   #1 = Methodref          #2.#3          // java/lang/Object."<init>":()V
14   #2 = Class               #4             // java/lang/Object
15   #3 = NameAndType         #5:#6          // "<init>":()V
16   #4 = Utf8                java/lang/Object
17   #5 = Utf8                <init>
18   #6 = Utf8                ()V
19   #7 = Fieldref            #8.#9          // java/lang/System.out:Ljava/io/PrintStream;
20   #8 = Class               #10            // java/lang/System
21   #9 = NameAndType         #11:#12        // out:Ljava/io/PrintStream;
22  #10 = Utf8                java/lang/System
23  #11 = Utf8                out
24  #12 = Utf8                Ljava/io/PrintStream;
25  #13 = String              #14            // Hello, world!
26  #14 = Utf8                Hello, world!
27  #15 = Methodref           #16.#17        // java/io/PrintStream.println:(Ljava/lang/String;)V
28  #16 = Class               #18            // java/io/PrintStream
29  #17 = NameAndType         #19:#20        // println:(Ljava/lang/String;)V
30  #18 = Utf8                java/io/PrintStream
31  #19 = Utf8                println
32  #20 = Utf8                (Ljava/lang/String;)V
33  #21 = Class               #22            // Main
34  #22 = Utf8                Main
35  #23 = Utf8                Code
36  #24 = Utf8                LineNumberTable
37  #25 = Utf8                LocalVariableTable
38  #26 = Utf8                this
39  #27 = Utf8                LMain;
40  #28 = Utf8                foo
41  #29 = Utf8                ()I
42  #30 = Utf8                SourceFile
43  #31 = Utf8                Main.java
```

Constructor

```
44 {
45   public Main();
46   descriptor: ()V
47   flags: (0x0001) ACC_PUBLIC
48   Code:
49     stack=1, locals=1, args_size=1
50     start local 0 // Main this
51     0: aload_0
52     1: invokespecial #1           // Method java/lang/Object."<init>":()V
53     4: return
54   end local 0 // Main this
55  LineNumberTable:
56     line 2: 0
57   LocalVariableTable:
58     Start   Length  Slot  Name   Signature
59     0        5      0   this   LMain;
60
```

Method foo

```
61 public int foo();
62 descriptor: ()I
63 flags: (0x0001) ACC_PUBLIC
64 Code:
65     stack=2, locals=1, args_size=1
66     start local 0 // Main this
67         0: getstatic      #7          // Field java/lang/System.out:Ljava/io/PrintStream;
68         3: ldc             #13         // String Hello, world!
69         5: invokevirtual  #15         // Method java/io/PrintStream.println:(Ljava/lang/String;)V
70         8: bipush          42
71        10: ireturn
72     end local 0 // Main this
73    LineNumberTable:
74         line 4: 0
75         line 5: 8
76     LocalVariableTable:
77         Start Length Slot Name Signature
78             0    11     0  this  LMain;
```


Execution Model / Data

- Data storage
 - Locals (parameter, local variables)
 - Stack (no registers)
 - Objects with fields
 - Static fields
- Methods
 - Static methods
 - Virtual methods

Execution Model / Data

- The JVM deals with 32 bit values
 - long and double values occupy 2 slots
- Constant pool
 - Strings, including
 - string literals used in programs
 - (full) class names, field & method names
 - method & field (type) descriptors
 - big integers, and floating-point numbers
 - dynamically computed constant (new, mostly for optimization)

JVM Instructions

- Each instruction has a name
 - optionally some immediate byte operands
 - typically they reference the constant pool
 - or offsets (for jumps)
 - rarely: immediate (small) integer values (max short)
 - they can modify the stack (pop values, push values, reorder them)
- One-byte **opcode**: max 256 instructions!
 - Could do variable-length opcodes if needed.
 - Instruction set very stable
- <https://docs.oracle.com/javase/specs/jvms/se16/html/jvms-6.html>
- https://en.wikipedia.org/wiki/Java_bytecode_instruction_listings

Bytecode Instructions (all 200 of them!)

Load/Store

- Push (**load**) variables on the stack, **store** stack value in variables
- Pattern: (A/F/D/I/L)(LOAD/STORE)[_0/_1/_2/_3]
 - Address (Object), Float, Double, Integer, Long
 - e.g. ALOAD, ISTORE_0
 - 50 instructions!
 - Simple form: one parameter (one byte) to give the variable index.

Load/Store (extra)

- IINC
 - Parameters: two bytes
 - Adds the parameter value to the integer variable.
- WIDE
 - "modifier" to put in front of LOAD/STORE/IINC to access a wider range of local variables.
 - Use three bytes of variable space index instead of just one.
 - Preserve opcode space (will be a common theme)

Constants

- Push constants onto the stack (20 instructions)
 - ACONST_NULL
 - ICONST(_M1/_0/_1/_2/_3/_4/_5)
 - FCONST(_0/_1/_2)
 - (L/D)CONST(_0/_1)
 - LDC, LDC_W, LDC2_W
 - BIPUSH, SIPUSH
- In theory, you could do everything with LDC[2]_W
- But give extremely common operations their own opcode to reduce executable size!

Type Conversions

- Everything has to be explicit (e.g. `1 + 2L` in Java)
- `value` \rightarrow `value`
- `(I/D/F/L)2(I/D/F/L)` (excluding identity)
- `I2(B/C/S)` (truncation)

Arithmetic/Binary Operations

- value, value \rightarrow value
- (D / F / I / L)(ADD / DIV / MUL / SUB / NEG / REM)
- (I / L)(AND / OR / SHL / SHR / USHR / XOR)

Stack Manipulation

- DUP[2][_X1/_X2]
 - Duplicates a value.
 - DUP2: Duplicates two values / a double-sized value.
 - X1: value2, value1 \rightarrow value1, value2, value1
 - X2: value3, value2, value1 \rightarrow value1, value3, value2, value1
- POP[2]
- SWAP

Arrays

- (A/F/D/I/L/C/B/S)A(LOAD/STORE)
 - load: arrayref, index → value
store: arrayref, index, value →
 - BALOAD: bytes or booleans! (otherwise booleans are int)
 - Why not store booleans as bits?
- [A]NEWARRAY
- MULTIANEWARRAY
 - parameters: type (constant pool ref), dimensions
 - count1, count2, ... → array
- ARRAYLENGTH

Field Access

- GETFIELD, PUTFIELD
 - Argument: field reference (constant pool ref)
 - Itself made of a class, name, field descriptor (type)
 - objectref \rightarrow value / objectref, value \rightarrow _
- GETSTATIC, PUTSTATIC
 - same
 - _ \rightarrow value / value \rightarrow _

Comparisons

- (D/F)CMP(G/L)
 - $A > B \rightarrow 1$, $A < B \rightarrow -1$, $A == B \rightarrow 0$
 - $A \text{ or } B == \text{NaN?}$ $G \rightarrow 1$, $L \rightarrow -1$
- LCMP
 - No ICMP! Directly embedded into jumps (next slide).
- INSTANCEOF
 - Parameter: class reference (constant pool ref)
 - $\text{True} == 1$, $\text{False} == 0$!
- CHECKCAST
 - Not really a comparison, throws an exception.

Jumps & Conditionals

- Parameter is always an offset in bytes (2 bytes, 4 for _W).
- GOTO[_W]
- IF_ACMP(EQ/NE) (two stack operands)
- IF_ICMP(EQ/NE/GE/GT/LE/LT) (two stack operands)
- IF(EQ/NE/GE/GT/LE/LT) (single stack operand)
 - $EQ = 0$, $GT > 0$, $LT < 0$
- IF[NON]NULL

Object Creation

- NEW
 - Parameter: class reference
- Calling the constructor?
 - INVOKESPECIAL <init>
 - The bytecode verifier checks that a constructor is called.

Method Invocation

- Parameter is always a method reference (class, name, method descriptor (type))
- INVOKESTATIC
 - Static methods, no this on stack.
- INVOKEVIRTUAL
 - "Regular" virtual calls, this on stack.
 - Get the object's class (may be different from static type) and get the method implementation from its method table.
 - The specific method implementation is not known at compile time (and might change each time the code is called)!
- INVOKINTERFACE
 - Same, but for interface methods..
- INVOKESPECIAL
 - *Static* linking, but with this pointer.
 - private methods, constructors, super methods
- INVOKEDYNAMIC
 - next slide

Invokedynamic (1)

- Branding: "for dynamic languages!". Yes, but why?
- The wormhole instruction
 - First call goes to a bootstrap method that returns a `java.lang.invoke.CallSite`
 - You have to write this yourself.
 - Then the instruction is "patched"
 - Next calls go to the call site directly
 - Can be optimized just like a regular call!
 - If you save the `CallSite`, you can change the target method!

Invokedynamic (2)

- Lets you write "polymorphic inline caches"
 - Will be explained in Truffle talk and/or optimization lecture.
 - Basically: always use a method tailored to the types that have actually been seen (+ a condition to check if the values are supported)

Misc Instructions

- [A/F/D/I/L]RETURN
- ATHROW
- MONITOR(ENTER/EXIT)
- NOP
- LOOKUPSWITCH / TABLESWITCH

Constructor

```
44 {
45   public Main();
46   descriptor: ()V
47   flags: (0x0001) ACC_PUBLIC
48   Code:
49     stack=1, locals=1, args_size=1
50     start local 0 // Main this
51     0: aload_0
52     1: invokespecial #1           // Method java/lang/Object."<init>":()V
53     4: return
54   end local 0 // Main this
55  LineNumberTable:
56     line 2: 0
57   LocalVariableTable:
58     Start Length Slot Name Signature
59     0      5      0 this LMain;
60 }
```

Method foo

```
61 public int foo();
62 descriptor: ()I
63 flags: (0x0001) ACC_PUBLIC
64 Code:
65     stack=2, locals=1, args_size=1
66     start local 0 // Main this
67         0: getstatic      #7          // Field java/lang/System.out:Ljava/io/PrintStream;
68         3: ldc            #13         // String Hello, world!
69         5: invokevirtual #15         // Method java/io/PrintStream.println:(Ljava/lang/String;)V
70         8: bipush          42
71        10: ireturn
72     end local 0 // Main this
73    LineNumberTable:
74         line 4: 0
75         line 5: 8
76     LocalVariableTable:
77         Start Length Slot Name Signature
78             0    11     0  this  LMain;
```


JIT: Register Allocation

- Most instructions map straightforwardly to underlying machine code
- However, most CPU architecture use registers
 - Registers are explicit (vs in-memory-hierarchy caches)
 - Mapping the JVM stack to memory would be too expensive
 - Most machine code instruction must use registers!
 - Use some register allocation algorithm
 - Good intro with a Java flavor:
<https://chrisseaton.com/truffleruby/register-allocation/>

Differences with Java

- (Generic) Type Erasure
 - The JVM doesn't have generic types.
 - `List<T>` is only `List`
 - Added backward-compatibly in Java 1.5
 - Generic types are not *reified*
- No checked exceptions
 - All exceptions act like `RuntimeException` in the JVM
- These two concerns are semantic analysis concerns.

Loading & Verification

- Class files are *loaded* into the JVM when required.
 - The first time the class is being referred.
 - Class references (from the constant pool) are strings!
- After loading the file, the bytecode is verified.
 - Misc verifications: constructors called, final variables not re-written after the constructor was called, ...
 - ***Typestate***: any instruction reachable through multiple paths (e.g. first instruction after an if-else) must be reached with the stack in the same state (size & types)

Typestate

Any instruction reachable through multiple paths must be reached with the stack in the same state (size & types).

- Java knows the type of every local & every slot on the stack.
 - Very different from machine code!
- It's not possible to iterate an array to copy it on the stack. (Why?)
- The stack can't overflow if you don't recurse.

Conclusion

- JVM
 - stack-based computation
 - storage: stack + locals + constant pool + fields
 - built-in GC — at the cost of memory control
 - built-in JIT compiler
- JVM Bytecode
 - low-level stack operations, jumps
 - transfers between stack and locals, fields, constant pool
 - highly-abstracted polymorphic method calls

Next Time

Generating JVM Bytecode with ASM